

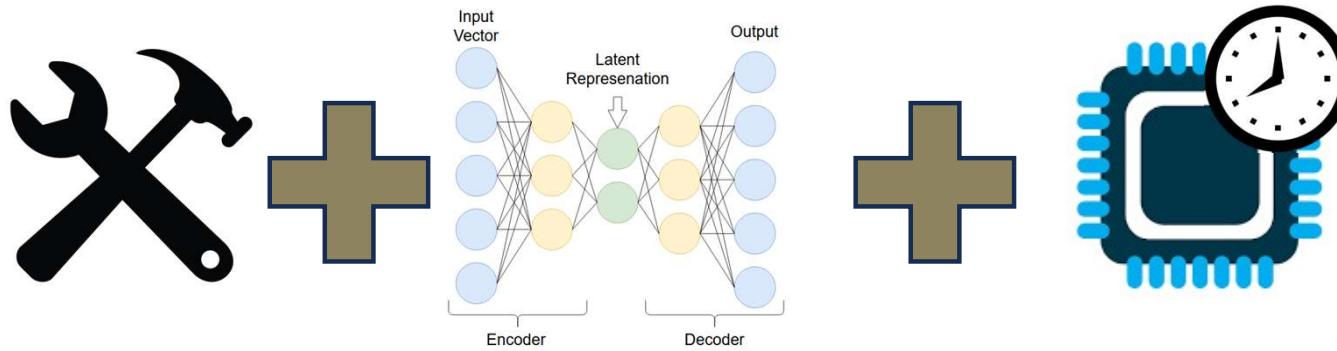
Machine Learning to Detect Attacks on CAN (Controller Area Network)



Finn Giegengack

Project Overview

- Deliverables
 - Machine learning pipeline to detect attacks on the CAN bus
 - Develop an autoencoder which can detect most attacks on CAN bus with reasonable accuracy



Project Overview

- Deliverables
 - Machine learning pipeline to detect attacks on the CAN bus
 - Develop an autoencoder which can detect most attacks on CAN bus with reasonable accuracy



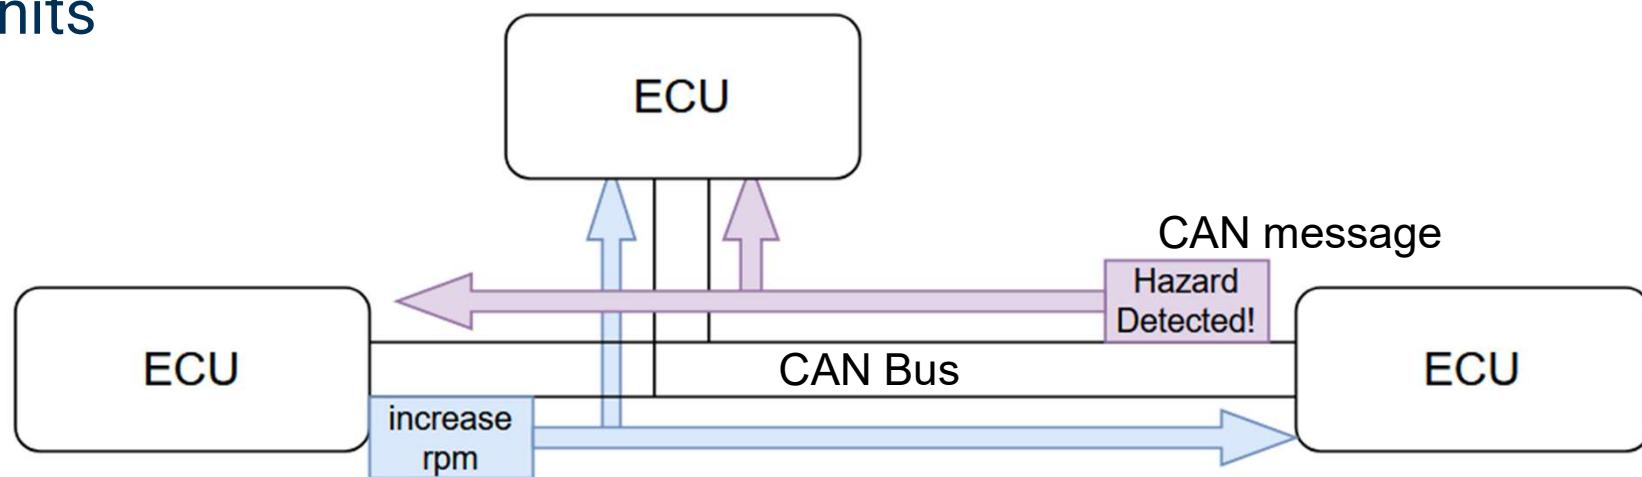
Background

Methods

Results

What's Next?

- CAN protocol
- Used commonly in automobiles
- Coordinates Communication between multiple Electronic Control Units



Background

Methods

Results

What's Next?

- Designed before wireless, over-the-air updates, ...
- **Threat model/assumptions have changed**

Q: What happens if a malicious actor gains access to bus?

Background

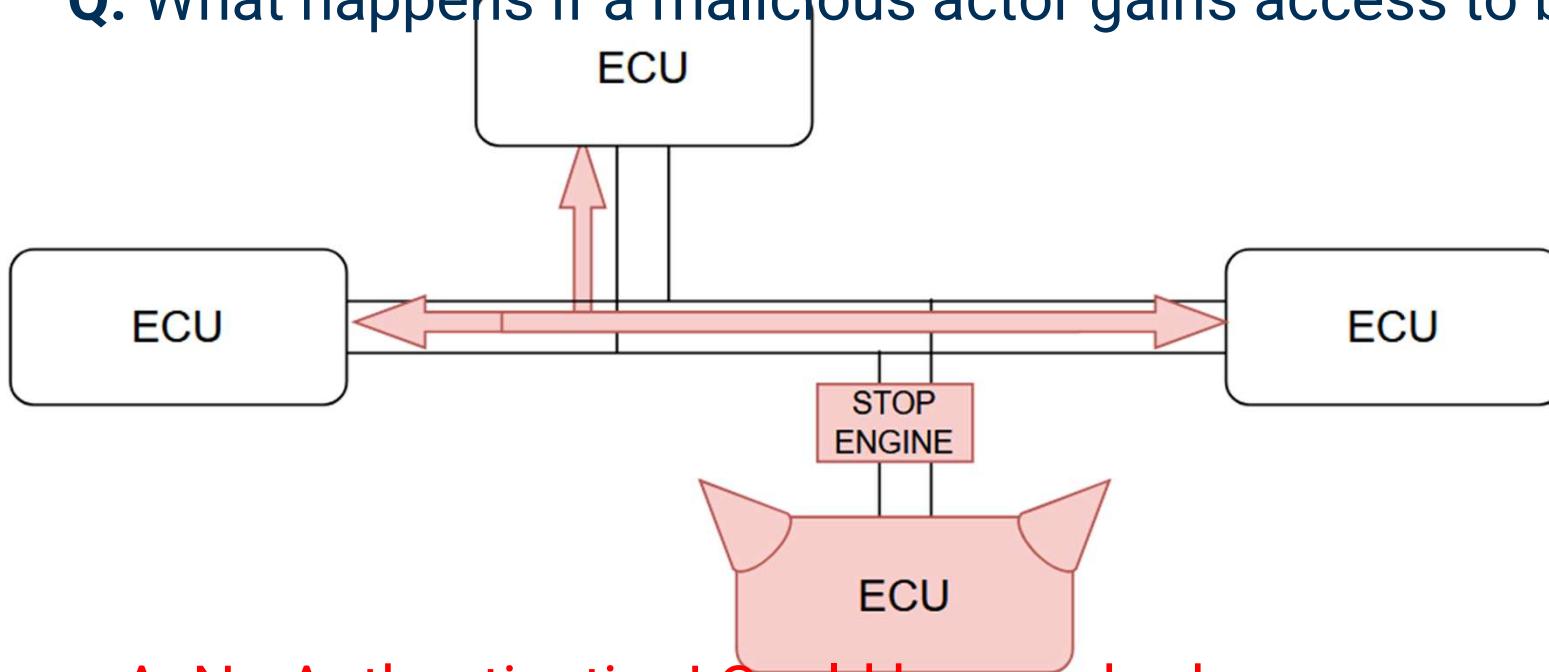
Methods

Results

What's Next?

- Designed before the internet, over-the-air updates, ...
- **Threat model/assumptions have changed**

Q: What happens if a malicious actor gains access to bus?



A: No Authentication! Could be very bad.

Background

Next?

TABLE I. SUMMARY OF EXPLOITS AGAINST THE CONTROLLER AREA NETWORK (CAN) BUS.

Vulnerability	Attack Type	Result
Unauthorized access	Packet injection. Reflashing ECU while driving.	Engine stopped. Code loaded into vehicles' telematics unit.
Unauthorized access Denial of Service	Packet injection to Body Control Module (BCM). Fuzzing.	Door lock relay activated. Wipers turned on/forced off. Trunk opened. Horn activated. Auto-headlight control deactivated. Use of washer fluid. Brake/auxiliary lights rendered inoperable.
Unauthorized access	Packet injection to Engine Control Module (ECM)	Engine timing and RPM disturbed. Engine cylinders stopped. Grind starter motor
Unauthorized access	Packet injection to Brake Control Module (BCM)	Brake application and release (evenly or unevenly) at speeds below 5 mph.
Denial of Service	Packet injection to other CAN-connected bus devices.	Disable CAN bus communication. Freeze instrument panel status.
Unauthorized access	Packet injection	Kill engine

[9] J. N. Brewer and G. Dimitoglou, "Evaluation of Attack Vectors and Risks in Automobiles and Road Infrastructure," in 2019 International Conference on Computational Science and Computational Intelligence (CSCI), Las Vegas, NV, USA: IEEE, Dec. 2019, pp. 84–89, ISBN 978-1-7281-5584-5. doi: 10.1109/CSCI49370.2019.00021. Accessed: Nov. 7, 2025. [Online].

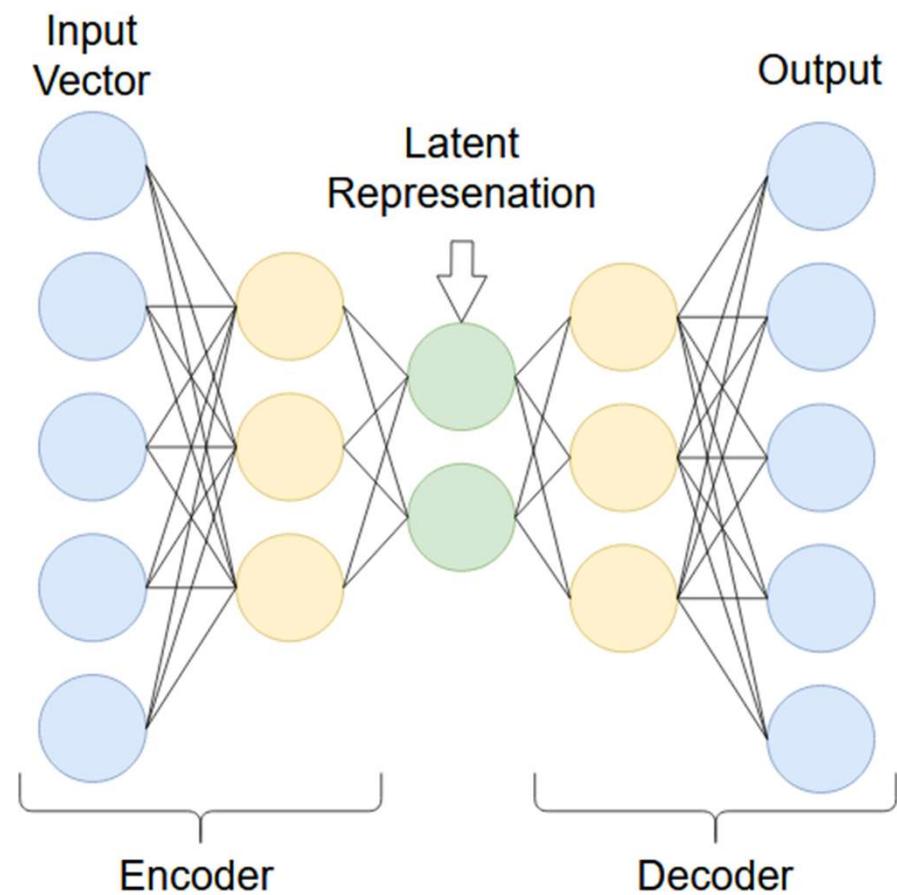
Background

Methods

Results

What's Next?

- What is an autoencoder?
- Learns to 'reconstruct' input



Background

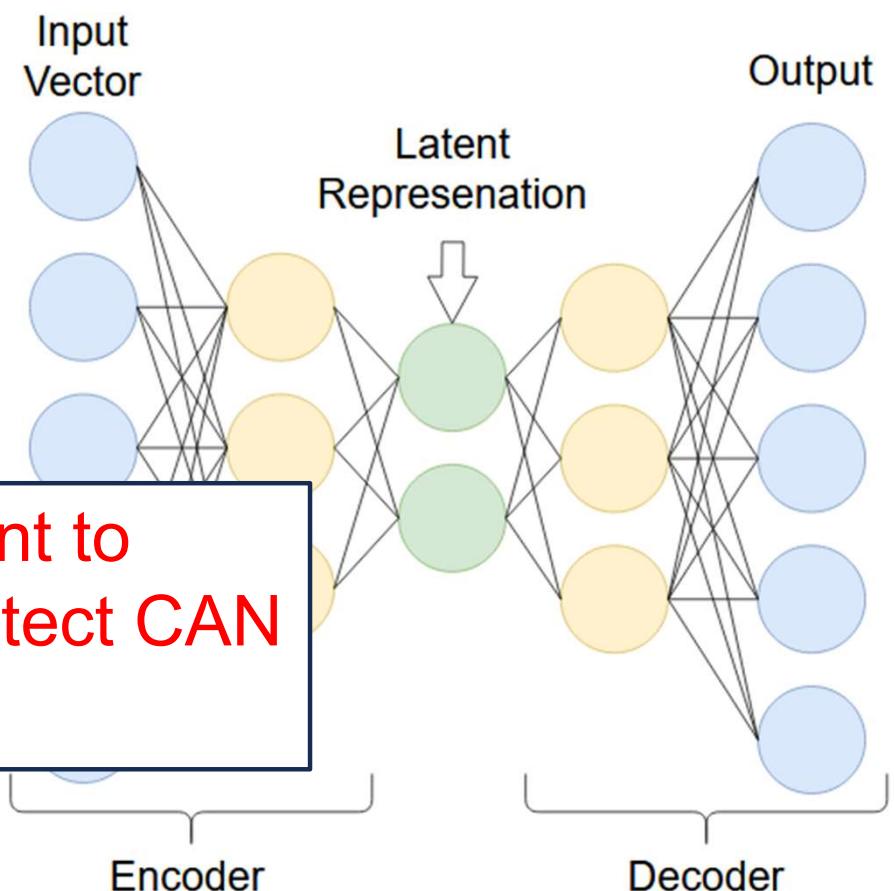
Methods

Results

What's Next?

- What is an autoencoder?
- Learns to 'reconstruct' input

But why would we want to reconstruct data to detect CAN bus attacks?



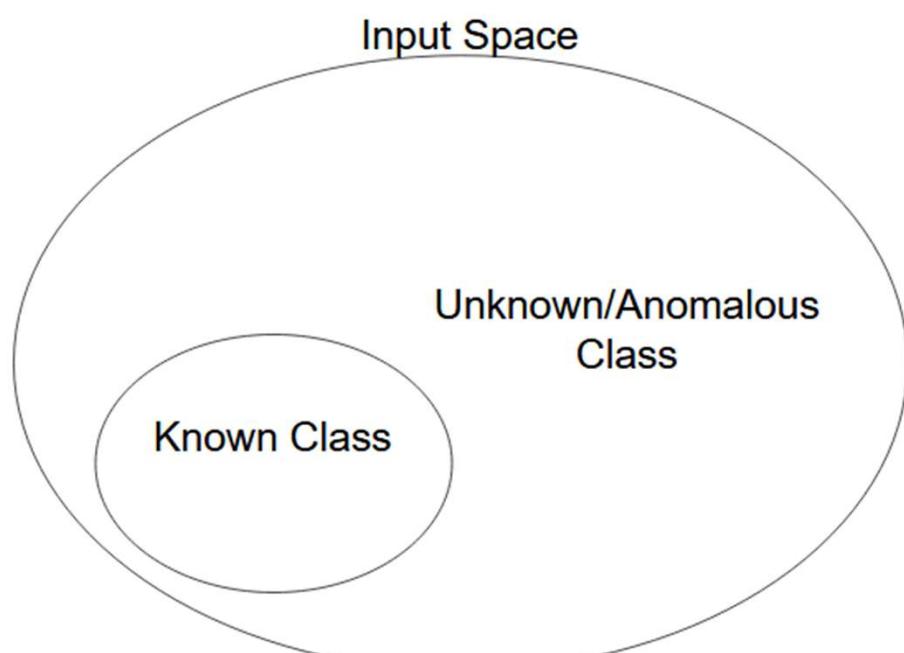
Background

Methods

Results

What's Next?

- One Class Classification
 - Subverts challenges with detecting CAN bus attacks: lack of attack data
 - We would like to train only on known data and detect anomalies



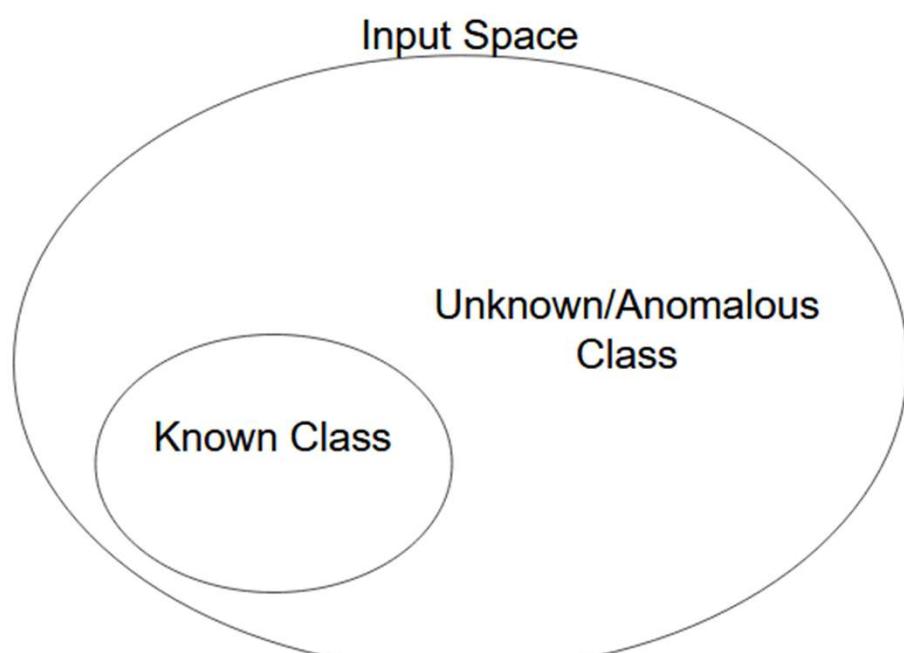
Background

Methods

Results

What's Next?

- One Class Classification
 - Subverts challenges with detecting CAN bus attacks: lack of attack data



- We would like to train only on known data and detect anomalies
- **Train only on Attack-Free CAN Bus DATA, detect attacks, which fall outside of known class**



Background

Methods

Results

What's Next?

- One Class Classification with Autoencoder
 - Train autoencoder on ‘good’ attack-free CAN bus data
 - Select reconstruction error threshold to maximize true positives, minimize false positives.
 - Challenges:
 - Autoencoder is characteristically good at generalizing to unseen inputs, but in this case that behavior is undesirable.
 - Possible Solutions:
 - Select latent dimension that is sufficiently small to promote overfitting on known class
 - **Use of Batch-Wise Feature Weighting block (BFW) as described in [6]**

[6] Y. Liao and B. Yang, “To Generalize or Not to Generalize: Towards Autoencoders in One-Class Classification,” in 2022 International Joint Conference on Neural Networks (IJCNN), Padua, Italy: IEEE, Jul. 18, 2022, pp. 1–8, isbn: 978-1-7281-8671-9. doi: 10.1109/IJCNN55064.2022.9892812. Accessed: Nov. 7, 2025. [Online].



Background

Methods

Results

What's Next?

- Dataset
 - can-train-and-test [8]
 - Contains attack-free CAN bus DATA and attack scenarios

```
autoencoder_can > data > 2016-chevrolet-silverado-DoS-labeled.csv
16240
16241
16242 timestamp,arbitration_id,data_field,attack
16243 1672531205.964396,000,0000000000000000,1
16244 1672531205.965094,1ED,6140020801B870DE,0
16245 1672531205.965409,000,0000000000000000,1
16246 1672531205.9661832,182,98C8010000000000,0
16247 1672531205.966197,2F9,1B010A0800003B,0
16248 1672531205.966221,000,0000000000000000,1
```

DoS Attack

8] B. Lampe and W. Meng, "Can-train-and-test: A New CAN Intrusion Detection Dataset," in 2023 IEEE 98th Vehicular Technology Conference (VTC2023-Fall), Hong Kong, Hong Kong: IEEE, Oct. 10, 2023, pp. 1–7, isbn: 979-8-3503-2928-5. doi: 10.1109/VTC2023-Fall60731.2023.10333756. Accessed: Nov. 7, 2025. [Online]. Available: <https://ieeexplore.ieee.org>.



Background

Methods

Results

What's Next?

- Dataset
 - can-train-and-test [8]
 - Contains attack-free CAN bus DATA and attack scenarios

dataset_type	total_messages	normal_messages	attack_messages	attack_percentage	duration_seconds	messages_per_second	unique_can_ids
attack-free	1254632	1254632	0	0	487.7283158	2572.399345	98
combined	920226	918266	1960	0.212991157	357.0358989	2577.404689	98
dos	252707	208924	43783	17.32559842	81.25668597	3109.984083	99
fuzzy	1059035	1008513	50522	4.770569433	392.2755001	2699.722516	2048
gear	1257602	1254632	2970	0.236163747	487.728138	2578.489741	98
interval	950657	943319	7338	0.771887232	366.811969	2591.673883	98
rpm	524212	522688	1524	0.290722074	203.2086239	2579.673982	98
speed	1255235	1254632	603	0.048038813	487.727381	2573.640622	98
standstill	524245	522688	1557	0.296998541	203.2089219	2579.832593	98

8] B. Lampe and W. Meng, "Can-train-and-test: A New CAN Intrusion Detection Dataset," in 2023 IEEE 98th Vehicular Technology Conference (VTC2023-Fall), Hong Kong, Hong Kong: IEEE, Oct. 10, 2023, pp. 1–7, isbn: 979-8-3503-2928-5. doi: 10.1109/VTC2023-Fall60731.2023.10333756. Accessed: Nov. 7, 2025. [Online]. Available: <https://ieeexplore.ieee.org>.

Background

Methods

Results

What's Next?

- Feature Extraction
 - Take next N messages
 - Timestamp for sample 1 = 0
 - Timestamp for subsequent messages, $n = \text{timestamp}_n - \text{timestamp}_1$
 - Pack timestamp, CAN ID, and (optionally) CAN Data broken up by byte for all N messages into a single input vector
- Consideration:
 - Ignore CAN Data Field
 - Other feature extraction methods described in [4] involve extracting density of certain messages in moving window.

[4] C. Chupong, N. Junhuathon, K. Kitwattana, T. Muankhaw, N. Ha-Upala, and M. Na-wong, “Intrusion Detection in CAN Bus using the Entropy of Data and One-class Classification,” in 2024 International Conference on Power, Energy and Innovations (ICPEI), Nakhon Ratchasima, Thailand: IEEE, Oct. 16, 2024, pp. 157–160, isbn: 979-83503-5677-9.doi: 10 . 1109 / ICPEI61831 . 2024 . 10748816. Accessed: Nov. 7, 2025.

Background

Methods

Results

What's Next?

- Evaluation
 - can-train-and-test includes 8 attack scenarios to test against
 - DoS, fuzzing
 - gear, interval, rpm, speed, standstill
 - involve spamming CAN message with known malicious effect, e.g. increasing rpm

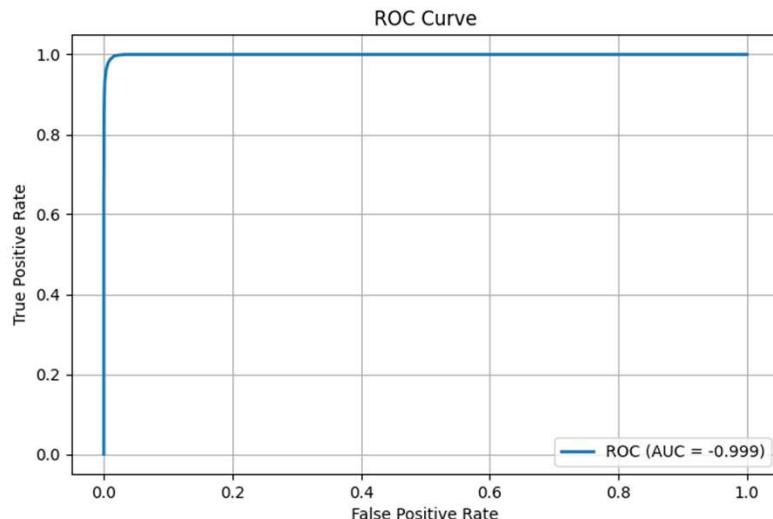
Background

Methods

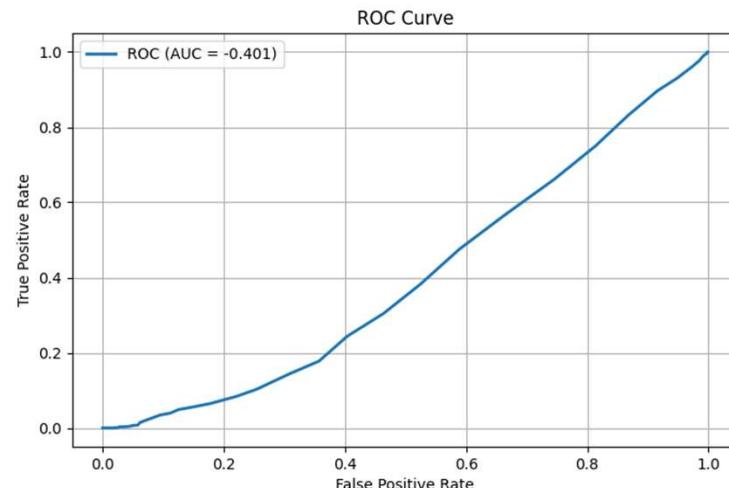
Results

What's Next?

- Analysis – ROC (Receiver Operating Characteristic) curve and AUC (Area Under Curve)
 - For one class classification, want to maximize the AUC of the ROC
 - ROC: plots true positive rate vs false positive rate for different threshold selections



Outstanding



Terrible



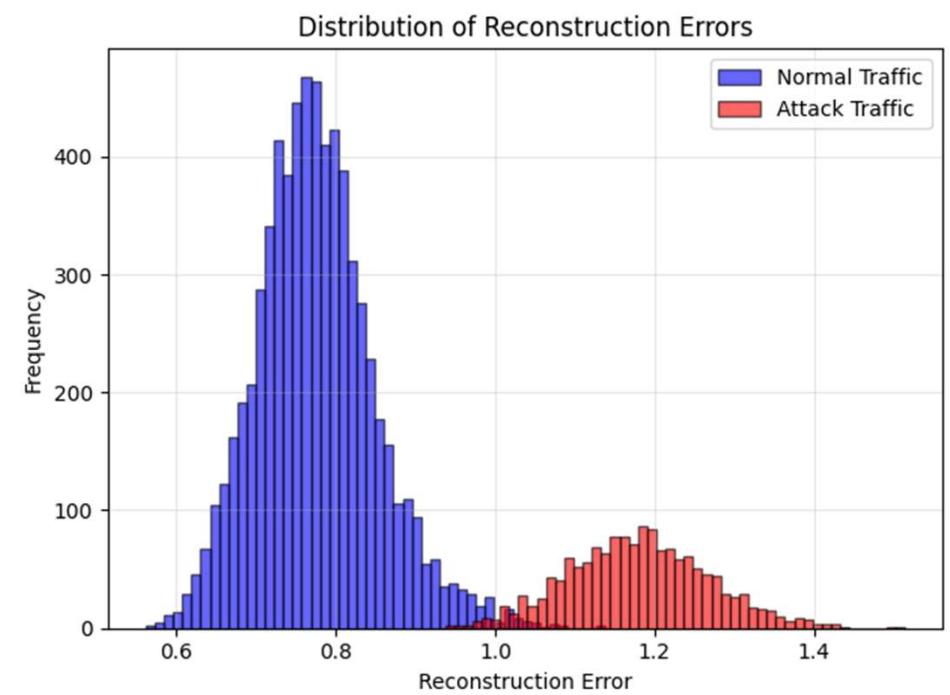
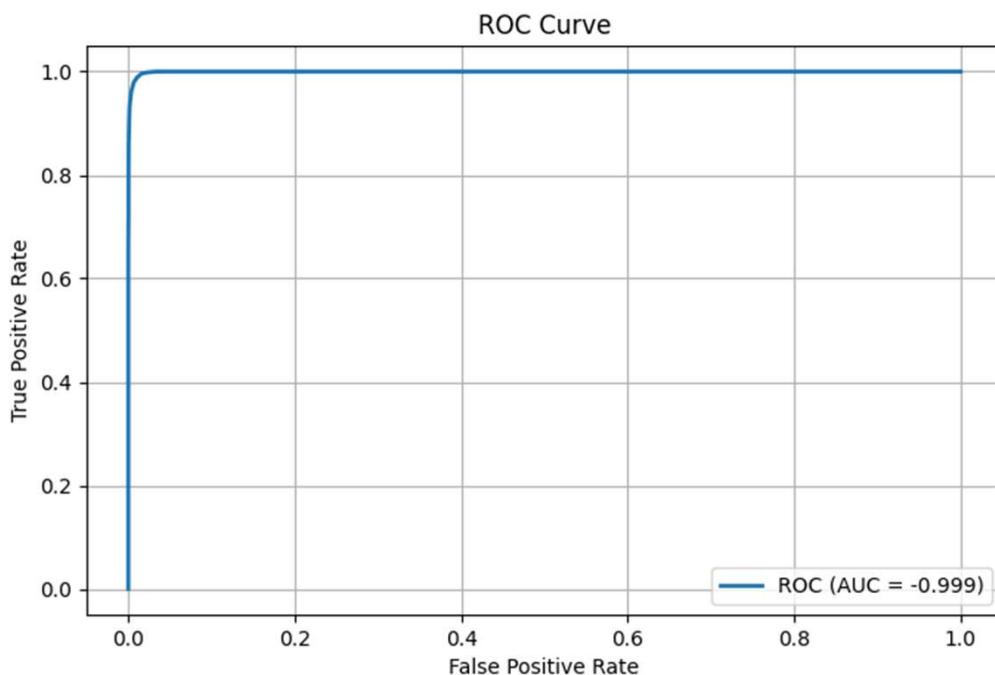
Background

Methods

Results

What's Next?

- Analysis – ROC Curve, another look



Background

Methods

Results

What's Next?

- Finding Optimal Parameters
 - N: number of CAN messages in a single input vector (size of window)
 - Learning rate
 - Batch size
 - Number of epochs
 - Feature extraction: whether to include CAN Data

Background

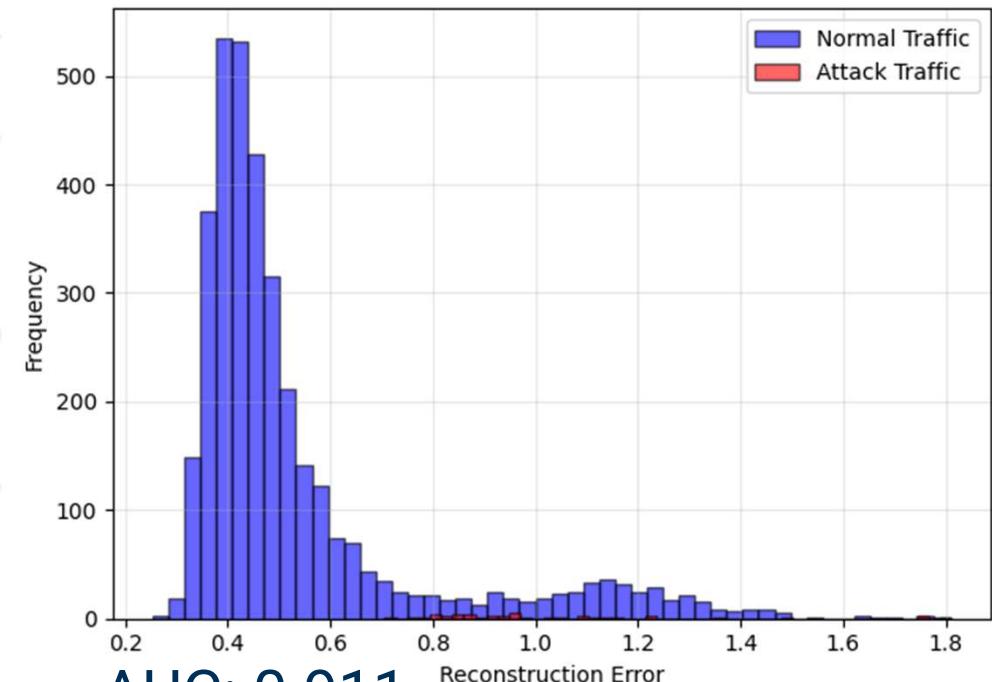
Methods

Results

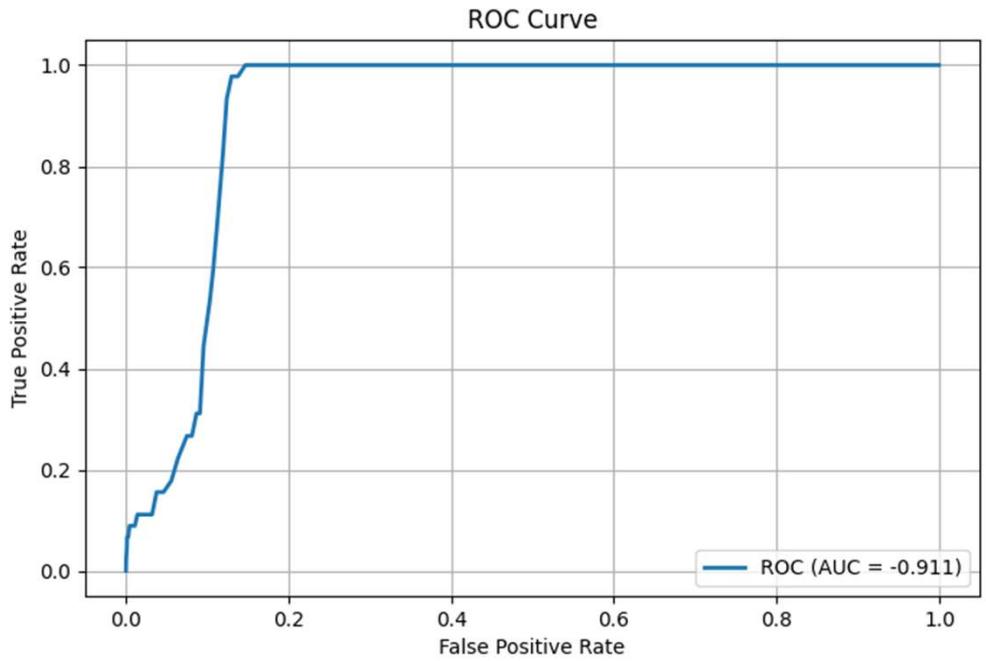
What's Next?

- Best results for Combined Attacks

Distribution of Reconstruction Errors



N=256, lr=0.0005, 2000 epochs, ignore CAN data



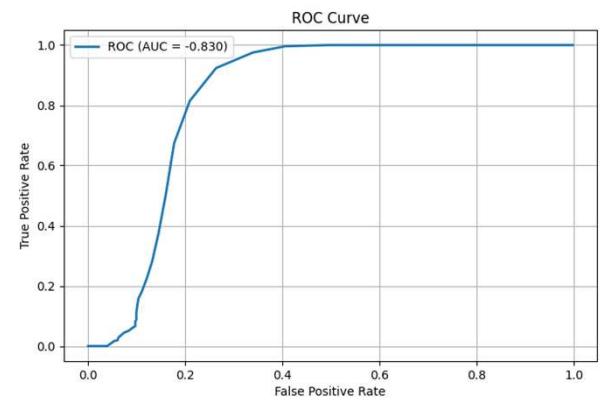
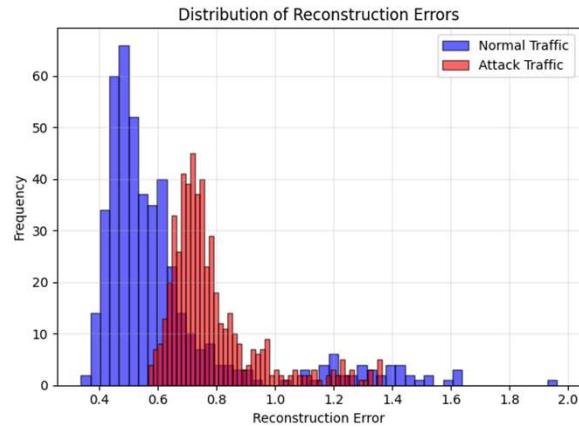
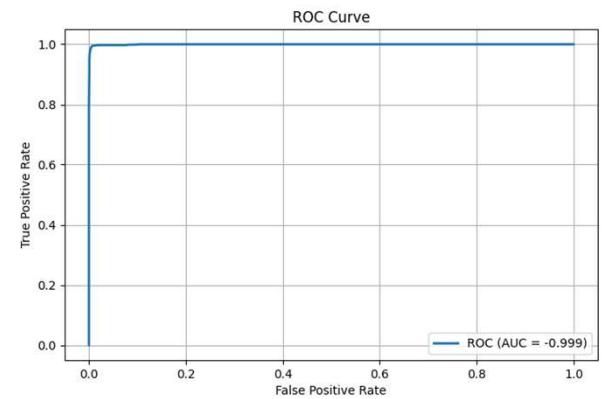
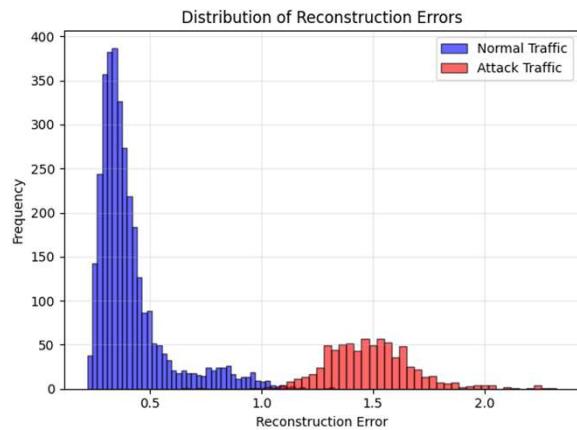
Background

Methods

Results

What's Next?

- Fuzzing Attacks



- DoS Attacks
 - Harder to predict?

Background

Methods

Results

What's Next?

- But!
- F1 Scores tell a different story

attack type	batch size	learning rate	epochs	msg per input	bfw	auc	max f1
combined	512	0.002	1000	256	1	0.932146771	0.188679245
dos	256	0.0018	1000	224	0	0.859448161	0.912710567
fuzzy	512	0.002	1000	224	0	0.999173709	0.998822144
gear	512	0.002	1000	256	1	0.90093119	0.215809285
interval	512	0.0001	1000	256	1	0.190671031	0.993494172
rpm	512	0.002	1000	256	1	0.906106613	0.275280899
speed	512	0.0005	1000	256	1	0.857480673	0.118012422
standstill	512	0.002	1000	256	1	0.906834394	0.277777778

Background

Methods

Results

What's Next?

- Effect of BFW on AUC

attack type	Average AUC w/ BFW	Average AUC w/o BFW	Difference in AUC mean	AUC Mean T Test p-value	Percent Improvement in AUC mean with BFW	in
combined	0.826095762	0.816539836	0.009555926	0.761548365	1.170295157	
dos	0.80550158	0.814371122	-0.008869542	0.714614287	-1.089127731	
fuzzy	0.977799241	0.977122082	0.000677159	0.97296619	0.069301324	
gear	0.683509097	0.657284032	0.026225065	0.477526173	3.989913629	
interval	0.572302369	0.564028981	0.008273388	0.84092411	1.466837379	
rpm	0.79943465	0.785840837	0.013593813	0.647115285	1.72984298	
speed	0.763801429	0.759404113	0.004397316	0.804667648	0.579048183	
standstill	0.797168415	0.786118394	0.011050021	0.702841116	1.40564338	

Background

Methods

Results

What's Next?

- Effect of BFW on F1 Score

attack type	Average F1 with BFW	Average F1 without BFW	Difference in F1 mean	F1 Mean T Test p-value	Percent Improvement in AUC mean with BFW
combined	0.117979391	0.111959004	0.006020388	0.590715696	5.37731434
dos	0.864217452	0.872967772	-0.00875032	0.591450952	-1.002364591
fuzzy	0.945836447	0.944790223	0.001046225	0.977353242	0.110736176
gear	0.086892499	0.078125988	0.008766511	0.378618931	11.22099173
interval	0.793008022	0.786842447	0.006165575	0.911550916	0.783584394
rpm	0.161738164	0.152070173	0.009667991	0.484399436	6.357585246
speed	0.075177555	0.073300329	0.001877225	0.689016663	2.561005362
standstill	0.162079824	0.155605887	0.006473938	0.644390019	4.160470943

Background

Methods

Results

What's Next?

- Try changing different parameters to improve F1
 - New feature extraction methods
 - Change latent dimension, dimensions of BFW
- Port model to Arduino Uno Q
- Explore other state of the art improvements to autoencoder for OCC



Georgia Tech College of Engineering
School of Electrical
and Computer Engineering



Thank You



Georgia Tech College of Engineering
School of Electrical
and Computer Engineering

References

- [1] F. Amato, L. Coppolino, F. Mercaldo, F. Moscato, R. Nardone, and A. Santone, "CAN-Bus Attack Detection With Deep Learning," *IEEE Transactions on Intelligent Transportation Systems*, vol. 22, no. 8, pp. 5081–5090, Aug. 2021, issn: 1524-9050, 1558-0016. doi: 10.1109/TITS.2020.3046974. Accessed: Nov. 7, 2025. [Online]. Available: <https://ieeexplore.ieee.org/document/9325944/>.
- [2] H. Ma, J. Cao, B. Mi, D. Huang, Y. Liu, and S. Li, "A GRU-Based Lightweight System for CAN Intrusion Detection in Real Time," *Security and Communication Networks*, vol. 2022, C. Chen, Ed., pp. 1–11, Jun. 27, 2022, issn: 1939-0122, 1939-0114. doi: 10.1155/2022/5827056. Accessed: Nov. 7, 2025. [Online]. Available: <https://www.hindawi.com/journals/scn/2022/5827056/>.
- [3] J. Guidry, F. Sohrab, R. Gottumukkala, S. Katragadda, and M. Gabbouj, "One-Class Classification for Intrusion Detection on Vehicular Networks," in *2023 IEEE Symposium Series on Computational Intelligence (SSCI)*, Mexico City, Mexico: IEEE, Dec. 5, 2023, pp. 1176–1182, isbn: 978-1-6654-3065-4. doi: 10.1109/SSCI52147.2023.10371899. Accessed: Nov. 7, 2025. [Online]. Available: <https://ieeexplore.ieee.org/document/10371899/>.

References (Contd.)

- [4] C. Chupong, N. Junhuathon, K. Kitwattana, T. Muankhaw, N. Ha-Upala, and M. Nawong, "Intrusion Detection in CAN Bus using the Entropy of Data and One-class Classification," in 2024 International Conference on Power, Energy and Innovations (ICPEI), Nakhon Ratchasima, Thailand: IEEE, Oct. 16, 2024, pp. 157–160, isbn: 979-8-3503-5677-9. doi: 10.1109/ICPEI61831.2024.10748816. Accessed: Nov. 7, 2025. [Online]. Available: <https://ieeexplore.ieee.org/document/10748816/>.
- [5] L. Liang et al., "Intrusion Detection Model for In-vehicle CAN Bus Based on TPE-LightGBM Algorithm," in 2025 IEEE 34th Wireless and Optical Communications Conference (WOCC), Taipa, Macao: IEEE, May 20, 2025, pp. 419–423, isbn: 979-8-3315-3928-3. doi: 10.1109/WOCC63563.2025.11082193. Accessed: Nov. 7, 2025. [Online]. Available: <https://ieeexplore.ieee.org/document/11082193/>.
- [6] Y. Liao and B. Yang, "To Generalize or Not to Generalize: Towards Autoencoders in One-Class Classification," in 2022 International Joint Conference on Neural Networks (IJCNN), Padua, Italy: IEEE, Jul. 18, 2022, pp. 1–8, isbn: 978-1-7281-8671-9. doi: 10.1109/IJCNN55064.2022.9892812. Accessed: Nov. 7, 2025. [Online]. Available: <https://ieeexplore.ieee.org/document/9892812/>.
- [7] B. M. T'oth and A. Bán'ati, "Autoencoder Based CAN BUS IDS System Architecture and Performance Evaluation," in 2025 IEEE 19th International Symposium on Applied Computational Intelligence and Informatics (SACI), Timisoara, Romania: IEEE, May 19, 2025, pp. 000 099–000 104, isbn: 979-8-3315-1547-8. doi: 10.1109/SACI66288.2025.11030168. Accessed: Nov. 7, 2025. [Online]. Available: <https://ieeexplore.ieee.org/document/11030168/>

References (Contd.)

- [8] B. Lampe and W. Meng, "Can-train-and-test: A New CAN Intrusion Detection Dataset," in 2023 IEEE 98th Vehicular Technology Conference (VTC2023-Fall), Hong Kong, Hong Kong: IEEE, Oct. 10, 2023, pp. 1–7, isbn: 979-8-3503-2928-5. doi: 10.1109/VTC2023-Fall60731. 2023.10333756. Accessed: Nov. 7, 2025. [Online]. Available: <https://ieeexplore.ieee.org/document/10333756/>.
- [9] J. N. Brewer and G. Dimitoglou, "Evaluation of Attack Vectors and Risks in Automobiles and Road Infrastructure," in 2019 International Conference on Computational Science and Computational Intelligence (CSCI), Las Vegas, NV, USA: IEEE, Dec. 2019, pp. 84–89, isbn: 978-1-7281-5584-5. doi: 10.1109/CSCI49370.2019.00021. Accessed: Nov. 7, 2025. [Online]. Available: <https://ieeexplore.ieee.org/document/9071073/>.
- [10] M. Kemmler, E. Rodner, E.-S. Wacker, and J. Denzler, "One-class classification with Gaussian processes," Pattern Recognition, vol. 46, no. 12, pp. 3507–3518, Dec. 2013, issn: 00313203. doi: 10.1016/j.patcog.2013.06.005. Accessed: Nov. 7, 2025. [Online]. Available: <https://linkinghub.elsevier.com/retrieve/pii/S0031320313002574>

References (Contd.)

- [11] J. Lee, S. Park, S. Shin, H. Im, J. Lee, and S. Lee, "ASIC Design for Real-Time CAN-Bus Intrusion Detection and Prevention System Using Random Forest," IEEE Access, vol. 13, pp. 129 856–129 869, 2025, issn: 2169-3536. doi: 10.1109/ACCESS.2025.3585956. Accessed: Nov. 7, 2025. [Online]. Available: <https://ieeexplore.ieee.org/document/11071663/>
- [12] F. Sohrab, J. Raitoharju, M. Gabbouj, and A. Iosifidis, "Subspace Support Vector Data Description," in 2018 24th International Conference on Pattern Recognition (ICPR), Beijing: IEEE, Aug. 2018, pp. 722–727, isbn: 978-1-5386-3788-3. doi: 10.1109/ICPR.2018.8545819. Accessed: Nov. 7, 2025. [Online]. Available: <https://ieeexplore.ieee.org/document/8545819/>.

Appendix A: CAN Bus Message Description

- Relevant Fields of CAN message

CAN ID: 11 bits

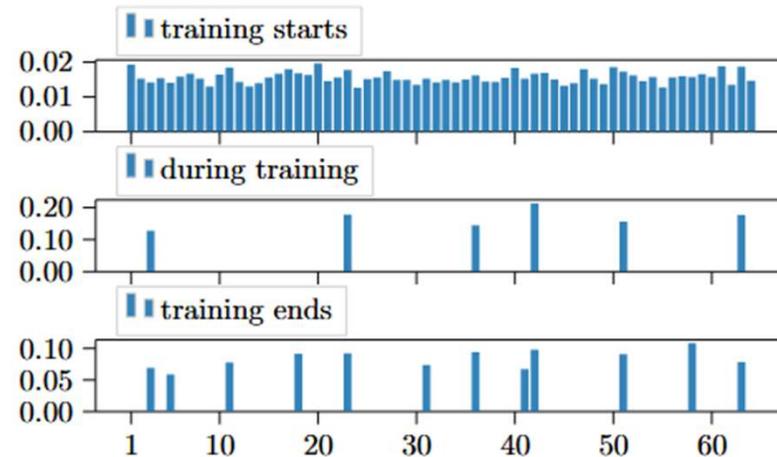
CAN Data: 0-8 bytes

- can-train-and-test gives us
 - CAN ID, CAN data
 - Timestamp (time in seconds)
 - Label (attack or non-attack)

8] B. Lampe and W. Meng, "Can-train-and-test: A New CAN Intrusion Detection Dataset," in 2023 IEEE 98th Vehicular Technology Conference (VTC2023-Fall), Hong Kong, Hong Kong: IEEE, Oct. 10, 2023, pp. 1–7, isbn: 979-8-3503-2928-5. doi: 10.1109/VTC2023-Fall60731.2023.10333756. Accessed: Nov. 7, 2025. [Online]. Available: <https://ieeexplore.ieee.org>.

Appendix B: Batch-Wise Feature Weighting

- Batch-Wise Feature Weighting (BFW)
 - Counteracts inherent ability of Autoencoder to generalize to new inputs
 - Assigns weighting to every value in latent dimension
 - Captures ‘most important’ latent features for known class, weighting others less



- Example of BFW training from Liao & Yang, 2022 [6]

[6] Y. Liao and B. Yang, “To Generalize or Not to Generalize: Towards Autoencoders in One-Class Classification,” in 2022 International Joint Conference on Neural Networks (IJCNN), Padua, Italy: IEEE, Jul. 18, 2022, pp. 1–8, isbn: 978-1-7281-8671-9. doi: 10.1109/IJCNN55064.2022.9892812. Accessed: Nov. 7, 2025. [Online].

Appendix B: Batch-Wise Feature Weighting

- Batch-Wise Feature Weighting (BFW)

```
# BFW: 2-layer network to compute w from z
self.W1 = nn.Linear(latent_dim, bfw_hidden_dim)
self.W2 = nn.Linear(bfw_hidden_dim, latent_dim)

self.b1 = nn.Parameter(torch.zeros(bfw_hidden_dim))
self.b2 = nn.Parameter(torch.zeros(latent_dim))
```

```
def forward(self, x):
    # Encode to latent representation z
    z = self.encoder(x)

    if self.use_bfw:
        inner_bfw = nn.functional.relu(self.W1(z) + self.b1)
        outer_bfw = self.W2(inner_bfw) + self.b2
        batch_mean = outer_bfw.mean(dim=0, keepdim=True)
        w = nn.functional.softmax(batch_mean, dim=1)
        z = z * w

    decoded = self.decoder(z)
    return decoded
```

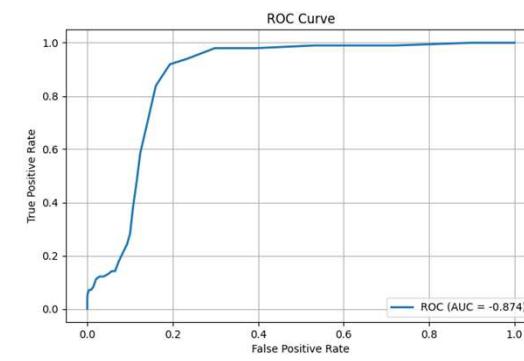
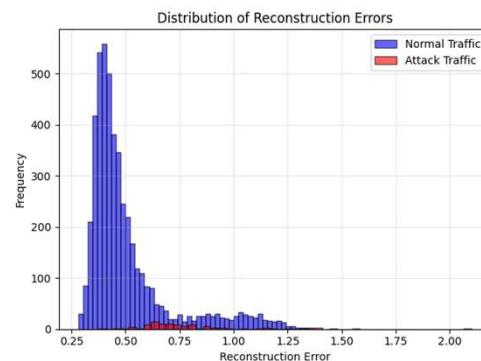
- Trainable Parameters

- Forward Pass applies BFW to latent representation

[6] Y. Liao and B. Yang, “To Generalize or Not to Generalize: Towards Autoencoders in One-Class Classification,” in 2022 International Joint Conference on Neural Networks (IJCNN), Padua, Italy: IEEE, Jul. 18, 2022, pp. 1–8, isbn: 978-1-7281-8671-9. doi: 10.1109/IJCNN55064.2022.9892812. Accessed: Nov. 7, 2025. [Online].

Appendix C: Results for Other Attacks

- Gear attack



- Rpm attack

